Mario Tapia-Pacheco

Los Angeles, CA | mtap1121@gmail.com | (323) 705-9808 | https://www.linkedin.com/in/mariotapiapacheco | https://github.com/mtapia-pacheco

Professional Summary

Data Science MS candidate at UC San Diego with experience building and evaluating machine learning pipelines, conducting statistical analysis, and translating quantitative findings into actionable insights. Skilled in Python, Scikit-learn, PyTorch, and SQL, with a track record of applying data science methods across real estate, public health, genomics, and NLP domains.

Education

University of California, San Diego
Master of Science in Data Science. GPA: 3.8/4.0. — *December 2026 (Expected)*

University of California, Santa Barbara
Bachelor of Science in Statistics and Data Science. GPA: 3.59/4.0 (Major GPA: 3.8/4.0). — *June 2024*

Experience

Data Science Intern — June 2026 – Present
IDX Exchange
- Building a machine learning pipeline with a 5-person team to predict California single-family property close prices using historical CRMLS sales data.
- Conducting preprocessing, feature engineering, and regression model evaluation using R^2, MAPE, and MdAPE.

Graduate-Level Evaluation Intern — November 2025 – April 2026
U.C. San Diego Center for Community Health
- Cleaned, restructured, and validated survey and program evaluation datasets for community health and nutrition incentive initiatives.
- Produced visual analyses and a stakeholder-facing evaluation report translating quantitative findings into program improvement recommendations.

Research Assistant, AI/ML for Genomics — March 2026 – Present
U.C. San Diego, The Amariuta Lab
- Contributed to EvoLen, a DNA foundation model research project studying tokenization strategies for genomic sequence modeling; manuscript currently in the rebuttal/discussion phase.
- Benchmarked EvoLen against models including Enformer, Nucleotide Transformer, and Grover across genomic prediction tasks; EvoLen ranks in the top 2 on 94.2% of non-splice tasks and 96.9% of regulatory/cross-species tasks (31/32).
- Organized experimental pipelines for dataset preparation, model evaluation, and results analysis in support of GenomicDiscoveryGym, an early-stage benchmark for evaluating AI scientist systems.

CHIRPS 3.0 Evaluation Researcher — January 2024 – August 2024
U.C. Santa Barbara Climate Hazards Center
- Applied statistical and time-series analysis to evaluate daily precipitation dataset performance across multiple CHIRPS 3.0 versions.

Projects

LLM-Powered Book Recommender System — January 2026
- Built an end-to-end semantic recommendation system using LangChain, Hugging Face, Chroma, vector search, and zero-shot classification over thousands of book descriptions.
- Optimized transformer inference and text-processing workflows via batching, truncation, and document chunking, reducing runtime from hours to minutes.

Skills

Programming: Python, R, SQL

ML/AI: Scikit-learn, PyTorch, TensorFlow/Keras, Hugging Face, LangChain, Chroma

Data/Tools: Pandas, NumPy, Matplotlib, Tableau, Power BI, Git/GitHub, AWS, Excel